


An abstract graphic on the left side of the slide featuring a dark blue background with a network of white and light blue nodes connected by thin white lines. The nodes are of varying sizes and are distributed across the left half of the image, creating a sense of depth and connectivity.

# Neural Networks to Distinguish Images

By Pedro Leite & Pedro  
Carvalho

# Convolution Neural Networks



1 Block VGG → 72%

2 Block VGG → 76%

3 Block VGG → 80%

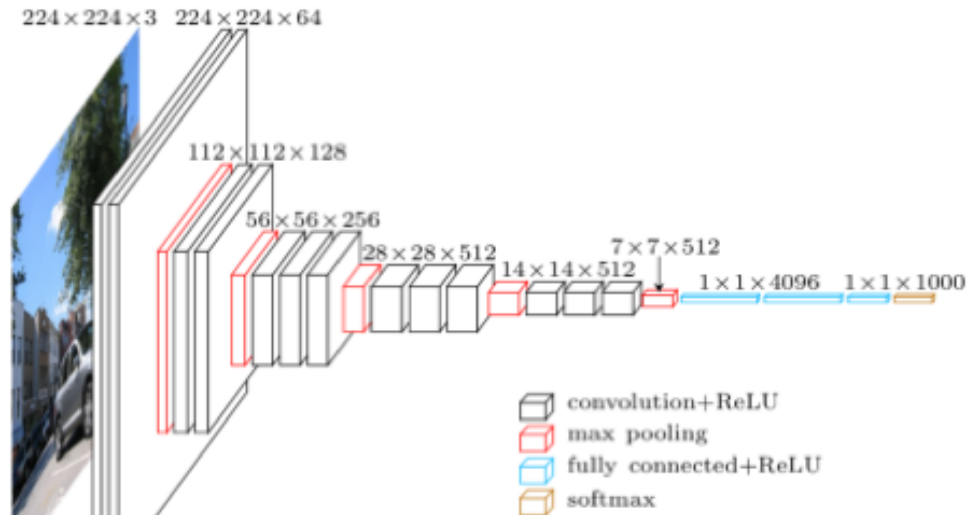
Overfitting

Dropout  
Regularization


Data  
Augmentation




# The VGG-16 Model



- 13 convulution layers.
- 3 connected layers.
- 1 max pooling layer.



ReLu



SoftMax


```
Epoch 1/5
293/293 [=====] - 2748s 9s/step - loss: 0.2493 - accuracy: 0.9644 - val_loss: 0.0751 - val_accuracy: 0.9770
Epoch 2/5
293/293 [=====] - 2747s 9s/step - loss: 0.0297 - accuracy: 0.9897 - val_loss: 0.0746 - val_accuracy: 0.9781
Epoch 3/5
293/293 [=====] - 2758s 9s/step - loss: 0.0095 - accuracy: 0.9976 - val_loss: 0.0806 - val_accuracy: 0.9779
Epoch 4/5
293/293 [=====] - 2749s 9s/step - loss: 0.0030 - accuracy: 0.9995 - val_loss: 0.0924 - val_accuracy: 0.9781
Epoch 5/5
293/293 [=====] - 2721s 9s/step - loss: 0.0013 - accuracy: 0.9998 - val_loss: 0.1008 - val_accuracy: 0.9781
C:\Users\35191\Desktop\TAIA\tester.py:66: UserWarning: "Model.evaluate_generator" is deprecated and will be removed in a future version. Please use "Model.evaluate", which supports generators.
  acc = model.evaluate_generator(test_it, steps=len(test_it), verbose=0)
> 97.811
```



# The Resnet Model

- 48 convulational layers.
- 1 max pooling layer.
- 1 average pooling layer.

```
Epoch 1/5  
625/625 [=====] - 1057s 2s/step - loss: 0.0996 - accuracy: 0.9650 - val_loss: 0.0754 - val_accuracy: 0.9696  
Epoch 2/5  
625/625 [=====] - 1057s 2s/step - loss: 0.0546 - accuracy: 0.9793 - val_loss: 0.0729 - val_accuracy: 0.9714  
Epoch 3/5  
625/625 [=====] - 1056s 2s/step - loss: 0.0428 - accuracy: 0.9833 - val_loss: 0.0806 - val_accuracy: 0.9680  
Epoch 4/5  
625/625 [=====] - 1058s 2s/step - loss: 0.0356 - accuracy: 0.9872 - val_loss: 0.0934 - val_accuracy: 0.9658  
Epoch 5/5  
625/625 [=====] - 1056s 2s/step - loss: 0.0272 - accuracy: 0.9897 - val_loss: 0.0787 - val_accuracy: 0.9728
```




# Adding Noisy Pictures



```
Epoch 1/5 [=====] - 1057s 2s/step - loss: 0.0996 - accuracy: 0.9650 - val_loss: 0.0754 - val_accuracy: 0.9696
Epoch 2/5 [=====] - 1057s 2s/step - loss: 0.0546 - accuracy: 0.9793 - val_loss: 0.0729 - val_accuracy: 0.9714
Epoch 3/5 [=====] - 1056s 2s/step - loss: 0.0428 - accuracy: 0.9833 - val_loss: 0.0806 - val_accuracy: 0.9680
Epoch 4/5 [=====] - 1058s 2s/step - loss: 0.0356 - accuracy: 0.9872 - val_loss: 0.0934 - val_accuracy: 0.9658
Epoch 5/5 [=====] - 1056s 2s/step - loss: 0.0272 - accuracy: 0.9897 - val_loss: 0.0787 - val_accuracy: 0.9728
```

# Adding Another Class



```
Epoch 1/5 [=====] - 2641s 9s/step - loss: 0.0334 - accuracy: 0.9413 - val_loss: -1.6923 - val_accuracy: 0.9650
Epoch 2/5 [=====] - 2626s 9s/step - loss: -26.3232 - accuracy: 0.9559 - val_loss: -507.1599 - val_accuracy: 0.9515
Epoch 3/5 [=====] - 2627s 9s/step - loss: -6792.3960 - accuracy: 0.9465 - val_loss: -124492.6875 - val_accuracy: 0.9204
Epoch 4/5 [=====] - 2630s 9s/step - loss: -1962060.7500 - accuracy: 0.9553 - val_loss: -35747860.0000 - val_accuracy: 0.9526
Epoch 5/5 [=====] - 2628s 9s/step - loss: -467379904.0000 - accuracy: 0.9420 - val_loss: -10380189696.0000 - val_accuracy: 0.9385
c:\Users\35191\Desktop\TAIA\tester.py:66: UserWarning: 'Model.evaluate_generator' is deprecated and will be removed in a future version. Please use 'Model.evaluate', which supports generators.
  acc = model.evaluate_generator(test_it, steps=len(test_it), verbose=0)
> 93.851
```



# Using Image Description to Train the Model

	A	B	C	D	E	F	G	H	I	J	K	L
1	Id,Subject Focus,Eyes,Face,Near,Action,Accessory,Group,Collage,Human,Occlusion,Info,Blur,Pawpularity											
2	0007de18844b0d	bb5e1f607da0606e0	0,1,1,1,0,0,1,0,0,0,0,0,63									
3	0009c66b9439883ba2750fb825e1d7db	0,1,1,0,0,0,0,0,0,0,0,0,42										
4	0013fd999caf9a3efe1352ca1b0d937e	0,1,1,1,0,0,0,0,1,1,0,0,28										
5	0018df346ac9c1d8413fcc888ca8246	0,1,1,1,0,0,0,0,0,0,0,0,15										
6	001dc955e10590d3ca4673f034feef2	0,0,0,1,0,0,1,0,0,0,0,0,72										
7	001dd4f6fafb890610b1635f967ea081	0,0,1,0,0,0,0,0,0,0,0,0,1,74										
8	0023b8a3abc93c712edd6120867deb53	0,1,1,1,0,0,0,0,0,1,1,0,0,22										
9	0031d6a9ef7340f898c3e05f92c7bb04	0,1,1,0,0,0,0,1,1,0,0,1,0,35										
10	0042bc5bada6d1cf8951f8f9f0d399fa	0,1,1,1,0,0,0,0,0,0,0,0,0,53										
11	0049cb81313c94fa007286e9039af910	0,1,1,1,0,0,0,0,0,0,0,0,0,21										
12	005017716086b8d5e118dd9fe26459b1	0,1,1,1,0,0,0,0,0,0,0,0,0,28										
13	00524dbf2637a80cbc80f70d3ff59616	0,1,1,1,0,0,0,0,0,0,0,0,0,2										
14	00630b1262efe301cb15a3b2022ba744	0,1,1,1,0,0,0,0,0,0,0,0,0,18										
15	006483b96ca9c09b7afed3e3d3af539d	0,1,1,1,0,0,0,0,0,1,1,0,0,41										
16	00655425c10d4c082dd7eeb97fa4fb17	0,1,0,0,0,1,0,0,0,0,0,0,0,13										
17	0067aaaa500b530c76b9c91af34b4cb8	0,1,1,1,0,0,0,0,0,0,0,0,0,98										
18	006cda7fec46a527f9f627f4722a2304	0,0,0,1,0,0,0,0,0,0,0,0,1,34										
19	006fe962f5f7e2c5f527b2e27e28ed6d	0,1,1,0,0,0,0,1,1,1,0,0,69										
20	0075ec6503412f21cf65ac5f43d80440	0,0,1,1,0,0,0,0,0,0,0,0,0,53										
21	00768659c1c90409f81dcdcecbd270513	0,1,1,0,0,0,0,0,0,0,0,0,0,100										
22	007c6ea0ccfca08d4736e5f459ed274a	0,1,1,1,0,0,0,1,0,0,0,0,0,29										
23	00838b02764e8d37edde39ed18693427	0,0,0,1,0,0,1,0,0,0,0,0,0,27										
24	0084a12f432ef6e47e8b9ff04d8a32b4	0,1,1,1,0,0,0,0,0,0,0,0,0,28										
25	0085bfb9a7ebfd776cb804d8b456bb05	0,1,1,1,0,0,0,0,0,0,0,0,0,50										
26	009367718f506d87cd898ed0310a2392	0,1,1,1,0,0,0,0,0,0,0,0,0,36										
27	0095f81bab3b68a4f70e99f0fcec7b06	0,0,1,1,0,0,0,0,1,0,0,1,15										
28	00a1ae8867e0bb89f061679e1cf29e80	0,1,1,0,0,0,1,0,0,0,0,0,0,52										
29	00a1e0cf89ff89a8f32d42e9025f6b2	0,1,1,1,0,0,0,0,0,0,0,0,0,40										
30	00a6d367daed96f77be6c91db493dfed	0,1,1,0,0,0,0,0,0,0,0,0,0,29										
31	00b151a572c9aabedf8cfce0fa18be25	0,0,1,1,0,0,0,0,0,0,0,0,1,21										
32	00c7b596627da7425b624d6b16c88a25	0,1,1,1,0,0,0,0,0,0,0,0,0,34										
33	00d1cb2ec8b263ae076ff95cae513a88	0,1,1,1,0,0,0,0,0,0,0,0,0,31										



ID: 007de18844b0d

Subject Focus: 0

Eyes: 0

Face: 1

...

Popularity: 63%



# Reinforcement Learning

## Image Classification by Reinforcement Learning with Two-State Q-Learning

Abdul Mueed Hafiz<sup>1</sup>

<sup>1</sup> Department of Electronics and Communication Engineering  
Institute of Technology, University of Kashmir  
Srinagar, J&K, India, 190006

ORC-ID: 0000-0002-2266-3708

### Reinforcement Learning for Visual Object Detection

Publisher: IEEE

[Cite This](#)

[PDF](#)

Stefan Mathe ; Aleksis Pirinen ; Cristian Sminchisescu [All Authors](#)

86

Paper  
Citations

1

Patent  
Citation

1644

Full  
Text Views



#### Abstract

##### Document Sections

1. Introduction
2. Related Work
3. Problem Formulation
4. Experiments and Results
5. Conclusions

#### Abstract:

One of the most widely used strategies for visual object detection is based on exhaustive spatial hypothesis search. While methods like sliding windows have been successful and effective for many years, they are still brute-force, independent of the image content and the visual category being searched. In this paper we present principled sequential models that accumulate evidence collected at a small set of image locations in order to detect visual objects effectively. By formulating sequential search as reinforcement learning of the search policy (including the stopping condition), our fully trainable model can explicitly balance for each class, specifically, the conflicting goals of exploration - sampling more image regions for better accuracy -, and exploitation - stopping the search efficiently when sufficiently confident about the target's location. The methodology is general and applicable to any detector response function. We report encouraging results in the PASCAL VOC 2012 object detection test set showing that the proposed methodology achieves almost two orders of magnitude speed-up over sliding window methods.